

Call Center Intel

Reference Deck · Idea Title: Call Center Intel — Agentic AI for Sales & Service Conversations
Working prototype, running today on a Medicare tele-sales reference implementation

The Opportunity

Contact centres have every tool to sell and serve — but nothing watches the conversation itself.

Compliance is reviewed weeks after the words were said. QA covers only a sampled handful of calls. New agents ramp on live prospects. Missed calls quietly become lost leads. And AI spend arrives as one untraceable invoice.

Call Center Intel puts a real-time AI layer on every conversation — enforcing compliance in the moment, coaching while the agent is still talking, scoring every call, and attributing the cost of doing so.

Built on an Insurance & Healthcare application, but suited to any contact centre whose agents speak directly to prospects or members.

The Solution — Four Pillars, One Pipeline

One real-time pipeline over the live conversation. The AI listens, assists and measures; the licensed human advises and closes.

1

Agentic Copilot

Works symbiotically with the application, not as a chatbot pushing text back: 23 governed tools that read live business data, drive the screens, draft emails and record updates for one-click approval, and place calls from natural language.

2

Conversation Intelligence

Live transcription with two-layer compliance — deterministic rules in under a second, plus an AI layer that catches paraphrased violations — alongside emotion-shift detection and proactive coaching on every call.

3

Supervisor & Governance

Live monitoring with every alert pinned to the line that triggered it. Business users edit compliance rules, coaching rules and QA scoring criteria themselves — effective on the very next call, with no deployment.

4

Cost Transparency

Per-call and per-agent AI cost from real usage logs, split by feature and drillable to the transcript behind any figure — instead of one opaque vendor invoice at the end of the month.

Technical Architecture

AI & agent layer

LangGraph copilot agent with 23 governed tools • QA evaluator • Claude via OpenRouter for reasoning • Deepgram streaming speech-to-text on both call tracks • live compliance, emotion and coaching analysis over the running transcript

Data & retrieval

Qdrant vector database (5 collections, nomic-embed-text 768-d via Ollama) • hybrid keyword + vector retrieval • MongoDB for application and reference data • hot-reloaded, versioned rule and QA-criteria sets • audit and cost telemetry per AI interaction

Voice, edge & platform

Twilio programmable voice (PSTN inbound & outbound, bidirectional media streams, signed webhooks) • Cloudflare edge with named tunnel • Caddy reverse proxy • Node 20 / Express 5 (TypeScript) API, containerised and cloud-portable

Application & delivery

React 19 / Vite single-page app with Twilio Voice SDK — agent workspace, call panel and copilot • supervisor command centre with live alert feed and mid-call transcript backfill • no-code admin for compliance, coaching and QA scoring criteria • RBAC

How & Where It Will Be Deployed

Environment	Containerised and horizontally scalable — cloud-portable by design. Deployable to any major cloud, hosted independently per client, or on client premises where data residency requires it. The data and AI layers can also run natively on platforms the client already owns, such as Databricks or Snowflake, so the solution sits where their data already lives.
Integration	Every layer is distinct and swappable — this stack is a reference implementation, not a dependency. React could be Angular; MongoDB could be SQL Server or Postgres; the telephony and speech vendors can change without touching the AI layer. REST + SSE surfaces over a governed tool layer front the client's existing CRM, telephony and master data, so it retro-fits onto what they already run.
Security & compliance	No sensitive identifiers requested, by design. Disabled capabilities are withheld from the model rather than discouraged in a prompt. Mandatory automated-call disclosure. Every AI interaction fully audited; PII handling and retention configured per client regulation.
Rollout approach	The four pillars are modular — a client takes what the business needs and leaves the rest. Copilot alone, compliance monitoring alone, or the full platform; each runs without the others. The usual path is to pilot one campaign or queue in shadow mode, switch on the pillars that matter most, then scale across queues and geographies — with the client's own business users tuning rules and QA criteria at every stage.
Ownership & support	EXL Services owns build, deployment, run and support. The client owns the data and the infrastructure it runs on, along with monitoring, latency thresholds and the escalation matrix. Rule and QA-criteria content stays with the client's own business users, who maintain it without a release cycle.

Where We Will Deploy & Sell

Deployment reach — EXL already supports the tele-sales departments of clients across regulated industries; the solution can be deployed into any number of them, and into new pursuits, with no client-specific rebuild

Where EXL already operates

Tele-sales departments of existing clients — the fastest path, because the operation, the data and the compliance obligations are already understood

Health-plan tele-sales and service desks — the domain the prototype already models end to end

Account teams — to identify and confirm the first target department

New logos / pipeline

US health plans — Medicare Advantage & ACA tele-sales desks

Banking & insurance — tele-sales desks operating under conduct rules

Collections & outreach BPOs — contact-practice and disclosure obligations

Retail & telecom — offer and claims compliance, churn-save coaching

Partnerships

GTM Strategy

Positioning

The AI layer for regulated conversations — not a new contact centre, an add-on over the one the client already runs. EXL consults on, builds and runs systems for some of the largest enterprises in the world; this is a complete end-to-end implementation, not a concept deck, so it can be demonstrated as-is from the first meeting of a pursuit.

Sales & channel

Expansion first: open new lines of business with organisations whose core operation is the contact centre, use this as living proof of what AI can do for them, then implement across the rest of their operation. Lead through existing account teams and attach to renewals; the live demo is the primary sales asset.

Pricing

Per-agent-seat subscription for copilot + conversation intelligence • per-minute pricing for automated inbound handling • outcome-linked tier for QA and compliance coverage • bundled into managed-operations pricing on BPaaS deals. Per-call cost telemetry keeps margin predictable.

Launch & marketing

Internal launch at the prototype showcase → reference case study from the first pilot → webinars on real-time compliance for regulated sales → thought leadership on human-in-the-loop AI → demo presence at payer and contact-centre industry events.